

Medi-Chat: An AI-Powered Chatbot to Provide Medical Information

Dibyajyoti Jana (dibyajyoti.jana.fit.cse22@teamfuture.in)
Department of Computer Science and Engineering
Future Institute of Technology
West Bengal, India

Abstract :

This paper introduces Medi-Chat, an AI-powered chatbot designed to deliver medical information. Leveraging advanced frameworks such as LangChain and employing a Retrieval-Augmented Generation (RAG) architecture, Medi-chat is capable of addressing the complexities of natural language understanding in a healthcare context. The integration of LLM enhances its performance in processing user queries and retrieving relevant knowledge from medical literature stored in PDF format. By combining efficient document retrieval with large language model reasoning, Medi-Chat ensures precise, contextually relevant, and trustworthy responses. This research outlines the methodologies applied in the development of Medi-Chat and evaluates its effectiveness as a tool for disseminating healthcare information.

Introduction :

The integration of artificial intelligence (AI) into healthcare has driven a transformative shift in the delivery of medical services, with AI-powered medical chatbots emerging as a significant innovation. These systems leverage natural language processing (NLP) and advanced AI algorithms to provide users with a wide range of healthcare-related services, including access to medical information, diagnostic assistance, and treatment recommendations. Such developments address the increasing demand for accessible healthcare solutions, particularly in light of the global shortage of medical professionals and the growing complexity of patient needs.

Medical chatbots are designed to improve patient engagement by delivering timely, accurate, and personalized responses. Their ability to communicate through natural language, combined with advanced algorithms for interpreting medical queries, enables them to simulate interactions with healthcare providers. This not only enhances patient access to reliable medical information but also supports healthcare professionals by streamlining routine inquiries and administrative tasks.

As the field progresses, researchers are increasingly focusing on improving medical chatbot performance by enhancing diagnostic accuracy, contextual relevance, and user experience. Advances such as context-aware processing and deep learning are driving this progress. Building on these developments, this research seeks to overcome existing limitations by integrating cutting-edge technologies to create a more robust, adaptive, and reliable AI-driven healthcare assistant.

Methodology:

The proposed system, Medi-Chat, is developed using a modular architecture that integrates both frontend and backend components, advanced Natural Language Processing (NLP) techniques, and a retrieval-augmented framework for reliable response generation. The methodology can be described as follows:

1. Frontend Development

The user interface (UI) is designed using React.js with TailwindCSS for styling.

The interface supports a chat-based interaction model, ensuring a simple, responsive, and user-friendly experience across devices.

2. Backend Development

The backend is implemented using Flask, a lightweight Python web framework.

Core functionalities include API endpoints, request handling, and integration with the LLM and vector store.

3. Authentication and Logging

Flask-Login is used for session management and user authentication.

Bcrypt is employed for secure password hashing.

Flask-SQLAlchemy is used for database management, ensuring secure storage of user credentials and interaction logs.

4. Data Source and Preprocessing

The system leverages medical textbooks and verified health literature as the primary knowledge base.

Text data is preprocessed using sentence segmentation, tokenization, and normalization to prepare it for embedding.

5. Vector Embeddings and Storage

Sentence-Transformer models are used to generate dense vector embeddings of the medical text.

These embeddings are stored and indexed in Pinecone, a vector database, via LangChain, enabling fast and efficient semantic search.

6. Large Language Model (LLM) Integration

Gemini, a transformer-based Large Language Model, is employed for natural language understanding and response generation.

The model is connected with the vector store through a Retrieval-Augmented Generation (RAG) pipeline to ensure responses are both contextually relevant and factually accurate.

7. Response Generation and Safety

For each user query:

The retriever searches Pinecone for relevant embeddings from the medical knowledge base.

The retrieved context is passed to Gemini for answer generation.

A safety filter validates the response to minimize bias, misinformation, or unsafe recommendations.

Final responses are returned to the user with proper citations where applicable.




Figure 1 : Data Preprocessing & Embedding Workflow

Working Principle:

The end user types a question (e.g. “What are diabetes symptoms?”)

Chat UI (React) shows chat bubbles.

On “Send”, it makes a POST request to backend with {message: "...user text..."}.

API (Flask) receives the request.

Validates request format.

Prepares to route it through the pipeline.

Authentication: JWT token (so only valid users can query).

Logging: Stores query metadata (userID, timestamp, query) in DB.

Helps with analytics, debugging, usage tracking.

(Optional) Rate limiting to prevent abuse.

Retriever (Vector DB + KB)

Converts user query → embedding vector.

Searches in Vector DB (Pinecone).

Retrieves top-k relevant chunks of knowledge base (KB).

Returns:

```
{
  "contexts": [
    { "content": "... diabetes symptoms ...", "source": "who_guidelines.txt" },
    { "content": "... frequent urination, fatigue ...", "source": "cdc_diabetes.txt" }
  ]
}
```

LLM (Generation)

Backend builds a prompt template:

SYSTEM: You are a medical assistant. Do not give prescriptions. Always cite sources.

CONTEXT:

Diabetes symptoms: increased thirst, frequent urination, fatigue. (Source: who_guidelines.txt)

- Diabetes symptoms: blurred vision, unexplained weight loss. (Source: cdc_diabetes.txt)

USER: What are diabetes symptoms?

Sends it to LLM (Gemin).

LLM generates a coherent, grounded response.

Safety Filter

Runs additional checks before sending back to user:

Content filter: detect unsafe / disallowed output (self-harm, violence).

Medical guardrails: if user asks “Should I take X drug?”, bot should respond:

“I cannot provide prescriptions. Please consult a doctor.”

Emergency escalation: if user says “I have chest pain”, response should redirect to emergency care.

Response + Citations

Final structured response is built:

```
{
  "answer": "Common diabetes symptoms include increased thirst, frequent urination, fatigue, blurred vision, and unexplained weight loss.",
  "sources": ["who_guidelines.txt", "cdc_diabetes.txt"]
}
```

Backend → Frontend

Backend returns JSON → frontend displays bot message.




Figure 2 : Working Diagram

Limitations :

Despite the success of Medi-Chat, there are several limitations that warrant further research. One key limitation is the reliance on the quality and scope of the medical literature used for training. While the chatbot performs well on standard queries, it may struggle with rare medical conditions or newly emerging diseases due to gaps in its training data. Additionally, the system's ability to handle real-time user feedback is limited, and continuous updates are required to improve its response accuracy. The model is also dependent on the quality of the initial data sources, which can introduce biases if not properly curated. Lastly, the chatbot is currently available only in English, restricting its accessibility to non-English speaking users

Conclusion :

This project demonstrates the potential of integrating advanced AI technologies into healthcare through the development of Medi-Chat. By leveraging LangChain, Retrieval-Augmented Generation (RAG) architecture, and the Gemini large language model, Medi-Chat effectively bridges the gap between medical literature and end-users, providing accurate, reliable, and context-aware healthcare information. The system not only enhances patient accessibility to trustworthy medical knowledge but also supports healthcare professionals by streamlining routine queries and administrative tasks. While the results highlight the promise of AI-powered chatbots in improving healthcare delivery, they also underscore the need for continuous updates, validation, and refinement to ensure reliability across diverse medical scenarios. Overall, Medi-Chat represents a step toward more intelligent, adaptive, and user-centric healthcare solutions, contributing to the broader vision of AI-driven healthcare innovation.

Future Scope :

While Medi-Chat demonstrates significant potential, several avenues remain for future enhancement:

Integration with Healthcare Systems –

Connecting Medi-Chat with electronic health records (EHRs) and hospital databases to provide personalized and real-time patient support.

Multilingual Support –

Expanding the chatbot's capability to interact in multiple languages, ensuring accessibility for a broader population.

Voice-Based Interaction –

Incorporating speech recognition and voice assistants to make the system more user-friendly, particularly for elderly or visually impaired users.

Advanced Personalization –

Using patient history and contextual data to deliver highly tailored healthcare recommendations.

Regulatory Compliance & Security –

Strengthening compliance with healthcare regulations (e.g., HIPAA) and ensuring robust data privacy and security.

Integration with Wearables & IoT –

Leveraging real-time health monitoring devices (smartwatches, fitness trackers) to provide proactive health insights.

By pursuing these directions, Medi-Chat can evolve into a comprehensive healthcare assistant, further bridging the gap between patients and medical professionals while ensuring reliability, accessibility, and trustworthiness in healthcare delivery.

Acknowledgements :

I would like to express my sincere gratitude to Future Institute of Technology for providing an excellent academic environment, modern facilities, and continuous support, all of which played a vital role in the successful completion of this project.

I am deeply indebted to my mentor, Mr. Subham Chakraborty, Faculty at Future Institute of Technology, for his invaluable guidance, unwavering encouragement, and constructive feedback throughout the course of this work. His profound knowledge, dedication to academic excellence, and commitment to nurturing my learning and professional development have been a constant source of inspiration and motivation.

References :

- [1] Adith Sreeram A S & Pappuri Jithendra Sai (2023). An Effective Query System Using LLMs and LangChain. International Journal of Engineering and Technical Research.12.
- [2] Andrew Reyner, Wibowo Tjiptomongsoguno, Audrey Chen, Hubert Michael Sanyoto, Edy Irwansyah(B), and Bayu Kanigoro. (2020). Medical Chatbot Techniques: A Review. 10.1007/978-3-030-63322-6 28.
- [3] R Jegadeesan, Dava Srinivas, N Umapathi, G Karthick, N Venkateswaran (2023). Section AResearch paper Personal Healthcare Chatbot for Medical Suggestions Using Artificial Intelligence and Machine Learning Eur. 12. 6004-6012. 10.31838/ecb/2023.12.s3.670.
- [4] Med-Bot: An AI-Powered Assistant to Provide Accurate and Reliable Medical Information.Ahan Bhatt, Nandan Vaghela (arXiv, Nov 2024).
- [5] Real-world Deployment and Evaluation of PErioperative AI CHatbot (PEACH) -- a Large Language Model Chatbot for Perioperative Medicine. Yu He Ke et al. (arXiv, Dec 2024).